

MLF_CNN

May 1, 2024

```
[1]: import pandas as pd
import numpy as np
from sklearn.datasets import fetch_openml
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout
import matplotlib.pyplot as plt
from sklearn.model_selection import KFold
from tensorflow.keras import layers, models, utils
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
from tensorflow.keras.utils import to_categorical # Correctly importing
↳to_categorical
from ucimlrepo import fetch_ucirepo
```

```
[2]: # Fetch dataset
optical_recognition_of_handwritten_digits = fetch_ucirepo(id=80)

# Data as pandas dataframes
X = optical_recognition_of_handwritten_digits.data.features
y = optical_recognition_of_handwritten_digits.data.targets

# Metadata
print(optical_recognition_of_handwritten_digits.metadata)

# Variable information
print(optical_recognition_of_handwritten_digits.variables)

# Optionally, print some information about the data
print("Features shape:", X.shape)
print("Target shape:", y.shape)
print("Unique digits in target:", np.unique(y))
```

```
{'uci_id': 80, 'name': 'Optical Recognition of Handwritten Digits',
'repository_url': 'https://archive.ics.uci.edu/dataset/80/optical+recognition+of
+handwritten+digits', 'data_url':
'https://archive.ics.uci.edu/static/public/80/data.csv', 'abstract': 'Two
```

versions of this database available; see folder', 'area': 'Computer Science', 'tasks': ['Classification'], 'characteristics': ['Multivariate'], 'num_instances': 5620, 'num_features': 64, 'feature_types': ['Integer'], 'demographics': [], 'target_col': ['class'], 'index_col': None, 'has_missing_values': 'no', 'missing_values_symbol': None, 'year_of_dataset_creation': 1998, 'last_updated': 'Wed Aug 23 2023', 'dataset_doi': '10.24432/C50P49', 'creators': ['E. Alpaydin', 'C. Kaynak'], 'intro_paper': {'title': 'Methods of Combining Multiple Classifiers and Their Applications to Handwritten Digit Recognition', 'authors': 'C. Kaynak', 'published_in': 'MSc Thesis, Institute of Graduate Studies in Science and Engineering, Bogazici University', 'year': 1995, 'url': None, 'doi': None}, 'additional_info': {'summary': 'We used preprocessing programs made available by NIST to extract normalized bitmaps of handwritten digits from a preprinted form. From a total of 43 people, 30 contributed to the training set and different 13 to the test set. 32x32 bitmaps are divided into nonoverlapping blocks of 4x4 and the number of on pixels are counted in each block. This generates an input matrix of 8x8 where each element is an integer in the range 0..16. This reduces dimensionality and gives invariance to small distortions.\r\n\r\nFor info on NIST preprocessing routines, see M. D. Garris, J. L. Blue, G. T. Candela, D. L. Dimmick, J. Geist, P. J. Grother, S. A. Janet, and C. L. Wilson, NIST Form-Based Handprint Recognition System, NISTIR 5469, 1994.', 'purpose': None, 'funded_by': None, 'instances_represent': None, 'recommended_data_splits': None, 'sensitive_data': None, 'preprocessing_description': None, 'variable_info': 'All input attributes are integers in the range 0..16.\r\nThe last attribute is the class code 0..9', 'citation': None}}

	name	role	type	demographic	description	units	\
0	Attribute1	Feature	Integer	None	None	None	
1	Attribute2	Feature	Integer	None	None	None	
2	Attribute3	Feature	Integer	None	None	None	
3	Attribute4	Feature	Integer	None	None	None	
4	Attribute5	Feature	Integer	None	None	None	
..	
60	Attribute61	Feature	Integer	None	None	None	
61	Attribute62	Feature	Integer	None	None	None	
62	Attribute63	Feature	Integer	None	None	None	
63	Attribute64	Feature	Integer	None	None	None	
64	class	Target	Categorical	None	None	None	

	missing_values
0	no
1	no
2	no
3	no
4	no
..	...
60	no
61	no
62	no

```

63         no
64         no

[65 rows x 7 columns]
Features shape: (5620, 64)
Target shape: (5620, 1)
Unique digits in target: [0 1 2 3 4 5 6 7 8 9]

```

```

[3]: def create_model():
    model = models.Sequential([
        layers.Conv2D(32, kernel_size=(3, 3), activation='relu',
        ↪input_shape=(8, 8, 1)),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Dropout(0.25),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.Flatten(),
        layers.Dense(128, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
    return model

model = create_model()
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 6, 6, 32)	320
max_pooling2d (MaxPooling2D)	(None, 3, 3, 32)	0
dropout (Dropout)	(None, 3, 3, 32)	0
conv2d_1 (Conv2D)	(None, 1, 1, 64)	18496
flatten (Flatten)	(None, 64)	0
dense (Dense)	(None, 128)	8320
dropout_1 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290

```
=====
Total params: 28426 (111.04 KB)
Trainable params: 28426 (111.04 KB)
Non-trainable params: 0 (0.00 Byte)
-----
```

```
[4]: # Convert features and labels to numpy arrays (if they are not already)
X = np.array(X)
y = np.array(y)

# Reshape X to fit the model input shape: (n_samples, 8, 8, 1)
X = X.reshape(-1, 8, 8, 1)

# Normalize the feature data
X = X.astype('float32') / 16 # Assuming pixel values range from 0 to 16

# One-hot encode the target labels
y = to_categorical(y, num_classes=10)

[5]: # Initialize lists to store test labels and predictions
y_true = []
y_pred = []
all_train_loss = []
all_val_loss = []
all_train_acc = []
all_val_acc = []

num_folds = 5
kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)
fold_no = 1

for train, test in kf.split(X):
    print(f'Training fold {fold_no}...')
    □
    ↪ print('-----')
    print(f'Training for fold {fold_no} ...')

    # Fit data to model
    history = model.fit(X[train], y[train],
                        batch_size=32,
                        epochs=10,
                        verbose=1,
                        validation_data=(X[test], y[test]))

    # Append loss and accuracy
    all_train_loss.extend(history.history['loss'])
    all_val_loss.extend(history.history['val_loss'])
```

```

all_train_acc.extend(history.history['accuracy'])
all_val_acc.extend(history.history['val_accuracy'])

# Generate generalization metrics
scores = model.evaluate(X[test], y[test], verbose=0)
print(f'Score for fold {fold_no}: {model.metrics_names[0]} of {scores[0]};  

↳ {model.metrics_names[1]} of {scores[1]*100}%')

# Predict on the test data
predictions = model.predict(X[test], batch_size=32)
y_pred.extend(np.argmax(predictions, axis=1))
y_true.extend(np.argmax(y[test], axis=1))

fold_no += 1

```

Training fold 1...

Training for fold 1 ...

Epoch 1/10

141/141 [=====] - 1s 3ms/step - loss: 1.4540 -
accuracy: 0.5347 - val_loss: 0.3917 - val_accuracy: 0.9128

Epoch 2/10

141/141 [=====] - 0s 2ms/step - loss: 0.4386 -
accuracy: 0.8648 - val_loss: 0.2316 - val_accuracy: 0.9288

Epoch 3/10

141/141 [=====] - 0s 2ms/step - loss: 0.3113 -
accuracy: 0.9070 - val_loss: 0.1645 - val_accuracy: 0.9466

Epoch 4/10

141/141 [=====] - 0s 2ms/step - loss: 0.2383 -
accuracy: 0.9315 - val_loss: 0.1366 - val_accuracy: 0.9573

Epoch 5/10

141/141 [=====] - 0s 2ms/step - loss: 0.1842 -
accuracy: 0.9464 - val_loss: 0.1118 - val_accuracy: 0.9671

Epoch 6/10

141/141 [=====] - 0s 2ms/step - loss: 0.1718 -
accuracy: 0.9475 - val_loss: 0.1018 - val_accuracy: 0.9751

Epoch 7/10

141/141 [=====] - 0s 2ms/step - loss: 0.1382 -
accuracy: 0.9584 - val_loss: 0.0866 - val_accuracy: 0.9786

Epoch 8/10

141/141 [=====] - 0s 2ms/step - loss: 0.1198 -
accuracy: 0.9646 - val_loss: 0.0820 - val_accuracy: 0.9742

Epoch 9/10

141/141 [=====] - 0s 2ms/step - loss: 0.1105 -
accuracy: 0.9682 - val_loss: 0.0727 - val_accuracy: 0.9786

Epoch 10/10

141/141 [=====] - 0s 2ms/step - loss: 0.1015 -
accuracy: 0.9693 - val_loss: 0.0658 - val_accuracy: 0.9795

Score for fold 1: loss of 0.0657673254609108; accuracy of 97.95373678207397%
36/36 [=====] - 0s 777us/step

Training fold 2...

Training for fold 2 ...

Epoch 1/10

141/141 [=====] - 0s 2ms/step - loss: 0.0937 -
accuracy: 0.9693 - val_loss: 0.0420 - val_accuracy: 0.9867

Epoch 2/10

141/141 [=====] - 0s 2ms/step - loss: 0.1025 -
accuracy: 0.9691 - val_loss: 0.0396 - val_accuracy: 0.9893

Epoch 3/10

141/141 [=====] - 0s 2ms/step - loss: 0.0853 -
accuracy: 0.9751 - val_loss: 0.0402 - val_accuracy: 0.9893

Epoch 4/10

141/141 [=====] - 0s 2ms/step - loss: 0.0697 -
accuracy: 0.9782 - val_loss: 0.0475 - val_accuracy: 0.9840

Epoch 5/10

141/141 [=====] - 0s 2ms/step - loss: 0.0749 -
accuracy: 0.9780 - val_loss: 0.0431 - val_accuracy: 0.9831

Epoch 6/10

141/141 [=====] - 0s 2ms/step - loss: 0.0745 -
accuracy: 0.9760 - val_loss: 0.0404 - val_accuracy: 0.9858

Epoch 7/10

141/141 [=====] - 0s 2ms/step - loss: 0.0705 -
accuracy: 0.9789 - val_loss: 0.0356 - val_accuracy: 0.9893

Epoch 8/10

141/141 [=====] - 0s 2ms/step - loss: 0.0592 -
accuracy: 0.9827 - val_loss: 0.0394 - val_accuracy: 0.9867

Epoch 9/10

141/141 [=====] - 0s 2ms/step - loss: 0.0593 -
accuracy: 0.9798 - val_loss: 0.0459 - val_accuracy: 0.9831

Epoch 10/10

141/141 [=====] - 0s 2ms/step - loss: 0.0518 -
accuracy: 0.9849 - val_loss: 0.0519 - val_accuracy: 0.9813

Score for fold 2: loss of 0.05190756916999817; accuracy of 98.13167452812195%
36/36 [=====] - 0s 829us/step

Training fold 3...

Training for fold 3 ...

Epoch 1/10

141/141 [=====] - 0s 3ms/step - loss: 0.0681 -
accuracy: 0.9775 - val_loss: 0.0187 - val_accuracy: 0.9920

Epoch 2/10

141/141 [=====] - 0s 2ms/step - loss: 0.0597 -
accuracy: 0.9804 - val_loss: 0.0165 - val_accuracy: 0.9929

Epoch 3/10

141/141 [=====] - 0s 2ms/step - loss: 0.0551 -

```

accuracy: 0.9824 - val_loss: 0.0185 - val_accuracy: 0.9938
Epoch 4/10
141/141 [=====] - 0s 2ms/step - loss: 0.0464 -
accuracy: 0.9878 - val_loss: 0.0223 - val_accuracy: 0.9902
Epoch 5/10
141/141 [=====] - 0s 2ms/step - loss: 0.0485 -
accuracy: 0.9840 - val_loss: 0.0211 - val_accuracy: 0.9920
Epoch 6/10
141/141 [=====] - 0s 2ms/step - loss: 0.0452 -
accuracy: 0.9844 - val_loss: 0.0193 - val_accuracy: 0.9920
Epoch 7/10
141/141 [=====] - 0s 2ms/step - loss: 0.0423 -
accuracy: 0.9858 - val_loss: 0.0226 - val_accuracy: 0.9938
Epoch 8/10
141/141 [=====] - 0s 2ms/step - loss: 0.0402 -
accuracy: 0.9878 - val_loss: 0.0248 - val_accuracy: 0.9920
Epoch 9/10
141/141 [=====] - 0s 2ms/step - loss: 0.0459 -
accuracy: 0.9851 - val_loss: 0.0206 - val_accuracy: 0.9929
Epoch 10/10
141/141 [=====] - 0s 2ms/step - loss: 0.0334 -
accuracy: 0.9893 - val_loss: 0.0232 - val_accuracy: 0.9929
Score for fold 3: loss of 0.02319428138434887; accuracy of 99.28825497627258%
36/36 [=====] - 0s 798us/step
Training fold 4...

```

```

-----
Training for fold 4 ...
Epoch 1/10
141/141 [=====] - 0s 2ms/step - loss: 0.0391 -
accuracy: 0.9878 - val_loss: 0.0063 - val_accuracy: 0.9991
Epoch 2/10
141/141 [=====] - 0s 2ms/step - loss: 0.0358 -
accuracy: 0.9907 - val_loss: 0.0115 - val_accuracy: 0.9956
Epoch 3/10
141/141 [=====] - 0s 2ms/step - loss: 0.0426 -
accuracy: 0.9873 - val_loss: 0.0125 - val_accuracy: 0.9947
Epoch 4/10
141/141 [=====] - 0s 2ms/step - loss: 0.0341 -
accuracy: 0.9889 - val_loss: 0.0088 - val_accuracy: 0.9964
Epoch 5/10
141/141 [=====] - 0s 2ms/step - loss: 0.0342 -
accuracy: 0.9898 - val_loss: 0.0122 - val_accuracy: 0.9956
Epoch 6/10
141/141 [=====] - 0s 2ms/step - loss: 0.0291 -
accuracy: 0.9915 - val_loss: 0.0160 - val_accuracy: 0.9947
Epoch 7/10
141/141 [=====] - 0s 2ms/step - loss: 0.0279 -
accuracy: 0.9902 - val_loss: 0.0130 - val_accuracy: 0.9964

```

```

Epoch 8/10
141/141 [=====] - 0s 2ms/step - loss: 0.0298 -
accuracy: 0.9907 - val_loss: 0.0108 - val_accuracy: 0.9964
Epoch 9/10
141/141 [=====] - 0s 2ms/step - loss: 0.0298 -
accuracy: 0.9900 - val_loss: 0.0126 - val_accuracy: 0.9947
Epoch 10/10
141/141 [=====] - 0s 2ms/step - loss: 0.0227 -
accuracy: 0.9929 - val_loss: 0.0136 - val_accuracy: 0.9947
Score for fold 4: loss of 0.013570494949817657; accuracy of 99.46619272232056%
36/36 [=====] - 0s 799us/step
Training fold 5...
-----
Training for fold 5 ...
Epoch 1/10
141/141 [=====] - 0s 3ms/step - loss: 0.0278 -
accuracy: 0.9913 - val_loss: 0.0040 - val_accuracy: 0.9982
Epoch 2/10
141/141 [=====] - 0s 3ms/step - loss: 0.0207 -
accuracy: 0.9944 - val_loss: 0.0048 - val_accuracy: 0.9973
Epoch 3/10
141/141 [=====] - 0s 2ms/step - loss: 0.0309 -
accuracy: 0.9902 - val_loss: 0.0063 - val_accuracy: 0.9973
Epoch 4/10
141/141 [=====] - 0s 2ms/step - loss: 0.0286 -
accuracy: 0.9911 - val_loss: 0.0100 - val_accuracy: 0.9973
Epoch 5/10
141/141 [=====] - 0s 2ms/step - loss: 0.0285 -
accuracy: 0.9920 - val_loss: 0.0071 - val_accuracy: 0.9982
Epoch 6/10
141/141 [=====] - 0s 2ms/step - loss: 0.0235 -
accuracy: 0.9927 - val_loss: 0.0111 - val_accuracy: 0.9956
Epoch 7/10
141/141 [=====] - 0s 2ms/step - loss: 0.0245 -
accuracy: 0.9918 - val_loss: 0.0132 - val_accuracy: 0.9956
Epoch 8/10
141/141 [=====] - 0s 2ms/step - loss: 0.0217 -
accuracy: 0.9935 - val_loss: 0.0092 - val_accuracy: 0.9964
Epoch 9/10
141/141 [=====] - 0s 2ms/step - loss: 0.0274 -
accuracy: 0.9904 - val_loss: 0.0069 - val_accuracy: 0.9973
Epoch 10/10
141/141 [=====] - 0s 2ms/step - loss: 0.0181 -
accuracy: 0.9944 - val_loss: 0.0086 - val_accuracy: 0.9956
Score for fold 5: loss of 0.008615042082965374; accuracy of 99.5551586151123%
36/36 [=====] - 0s 794us/step

```



```
[6]: # Generate the confusion matrix
cm = confusion_matrix(y_true, y_pred)
print("Confusion Matrix:")
print(cm)

# Classification report
report = classification_report(y_true, y_pred)
print("Classification Report:")
print(report)

# Plotting the confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt="d", cmap='Blues')
plt.title('Confusion Matrix')
plt.ylabel('Actual Label')
plt.xlabel('Predicted Label')
plt.show()

# Calculate average loss and accuracy
average_loss = np.mean(all_val_loss)
average_accuracy = np.mean(all_val_acc) * 100

print(f'Average Loss: {average_loss}')
print(f'Average Accuracy: {average_accuracy}%')

# Plot training & validation loss values
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(all_train_loss, label='Training Loss')
plt.plot(all_val_loss, label='Validation Loss')
plt.title('Loss across all folds')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

# Plot training & validation accuracy values
plt.subplot(1, 2, 2)
plt.plot(all_train_acc, label='Training Accuracy')
plt.plot(all_val_acc, label='Validation Accuracy')
plt.title('Accuracy across all folds')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

Confusion Matrix:

```
[[552  0  0  0  1  0  1  0  0  0]]
```


```

[ 0 564  2  0  0  0  0  1  3  1]
[ 0  0 556  0  1  0  0  0  0  0]
[ 0  0  2 562  0  5  0  0  1  2]
[ 0  0  0  0 559  0  5  0  0  4]
[ 0  0  0  2  1 549  0  0  1  5]
[ 0  1  0  0  2  0 554  0  1  0]
[ 0  0  0  1  0  0  0 565  0  0]
[ 1  1  0  1  0  1  0  0 548  2]
[ 0  1  0  4  2  3  0  3  1 548]]

```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	554
1	0.99	0.99	0.99	571
2	0.99	1.00	1.00	557
3	0.99	0.98	0.98	572
4	0.99	0.98	0.99	568
5	0.98	0.98	0.98	558
6	0.99	0.99	0.99	558
7	0.99	1.00	1.00	566
8	0.99	0.99	0.99	554
9	0.98	0.98	0.98	562
accuracy			0.99	5620
macro avg	0.99	0.99	0.99	5620
weighted avg	0.99	0.99	0.99	5620



Average Loss: 0.04553227401338518
Average Accuracy: 98.61743795871735%

